

Federated Learning for Privacy-Preserving Healthcare Diagnostics

A Comprehensive Analysis of Distributed Machine Learning in Medical Data
Protection

Teams Members

1. Hasan Alaa Srour – 1320224116

2. Mohammed Emad Al-maswabi - 1320222450

Research Paper

Artificial Intelligence & Healthcare Informatics
Privacy-Preserving Machine Learning | Medical Diagnostics

2026

Introduction

Machine learning has changed how we find diseases early, like diabetes, by checking patient data. But health info is private, with rules like HIPAA in the US and GDPR in Europe. Usually, machine learning puts all data in one spot, which can cause privacy problems and make it hard for hospitals to work together. This is a big deal in places where healthcare is all over the place, because it stops them from creating good ways to find diseases.

Federated Learning is a new way to fix this. **(McMahan et al., 2017)** came up with it. It lets hospitals train models together without sharing patient data. Instead, they only share updates to the model, like changes to its settings, which are put together on a server. This keeps data safe and private.

For healthcare, Federated Learning lets hospitals train models together using their own data. This makes better models for things like predicting diabetes, all while keeping patient info safe **(Rieke et al., 2020)**. It helps when hospitals can't share data because of rules or ethics. They can still learn from each other without risking privacy.

This research looks at the improvements in Federated Learning, checks how it works compared to older methods using the PIMA Indians Diabetes dataset, looks at examples in healthcare, talks about ethical stuff like privacy and bias, and suggests where research can go to make AI in healthcare more private.

Literature Review

Federated Learning was first put together by **(McMahan et al., 2017)**. It's about learning from different sources without moving the data. The main idea, called Federated Averaging (FedAvg), is that everyone shares updates to the model but keeps their own data. This saves time compared to sharing data all the time. It started to help with mobile apps and keeping consumer data private.

Later work added better privacy to Federated Learning. **(Geyer et al., 2018)** added noise to the model updates to stop people from figuring out the training data. This showed that you can balance privacy and how well the model works. Others made ways to keep the server from seeing individual updates and changed how the model learns to deal with different data, which is common in healthcare because patients and practices vary across hospitals **(Tahir et al., 2025)**.

In healthcare, Federated Learning has been used for records and images. **(Rieke et al., 2020)** said it could help hospitals work together while following privacy rules. Studies have used it to predict diabetes with datasets like the PIMA Indians Diabetes dataset, getting about 95% to 99% of the accuracy of normal methods while keeping privacy **(Kumar et al., 2024)**. Reviews say it helps with finding disease risks across hospitals, but it still has problems like slow communication, different data, and attacks **(Tahir et al., 2025)**.

Analysis

Federated Learning is way better than the old way because it lets hospitals train models together without sharing private data. This fixes the privacy issue in healthcare and still keeps the models working well.

To see how well it works, we used the PIMA Indians Diabetes dataset. It has 768 patients with info like glucose levels, BMI, age, and family history to predict diabetes. We used a simple neural network for both normal and federated training.


In the normal setup, we put all the data together and trained the model, getting 77.27% accuracy. This is the best you can do when you don't worry about privacy.

For federated learning, we split the data into five groups to act like different hospitals with different patients. We used the Federated Averaging way for 20 rounds. Each group updated the model with their own data and sent the updates to the server. This way, the model got 75.32% accuracy, which is about 97.5% of the normal way. This matches what others have found, where Federated Learning keeps about 95% to 99% of the accuracy (**Kumar et al., 2024; McMahan et al., 2017; Tahir et al., 2025**).

Here's a table to compare:

Approach	Test Accuracy	Accuracy Retention	Privacy Level	Data Sharing Required
Centralized	77.27%	100%	Low (full data exposure)	Yes
Federated (FedAvg)	75.32%	97.5%	High (local data only)	No

This shows that Federated Learning is a good balance between privacy and how well the model works. It loses a little accuracy compared to normal training, but it protects privacy, follows rules, and lowers the risk of data leaks. It also lets hospitals work together to make better models without breaking privacy rules.



Applications

Federated Learning is used in healthcare, letting hospitals work together on AI while keeping patient data private and following rules.

One big area is medical images, where Federated Learning helps make models without putting all the images in one place. NVIDIA Clara and the German RACOON use it to train models on different datasets to find tumors, check breast density, and look at heart images. This helps hospitals improve find rare diseases without risking privacy (**Rieke et al., 2020**).

For predicting diseases like diabetes, Federated Learning helps make risk models using health records from different clinics. This helps with early treatment while keeping data safe. In cancer, companies like Owkin use federated networks to help find drugs and predict how people will respond to treatment without seeing patient records.

It's also used for watching diseases and modeling outbreaks, like with COVID-19. And with wearable devices, Federated Learning can use data from these to update models while keeping user data private.

Who's using it:

- Hospitals and research groups, like Columbia University, that use it for trials and studies.
- Drug companies that use it to research and predict results while keeping privacy.
- Health groups that use it to watch diseases and manage health across big groups of people.

Here's a table of use cases:

Use Case	Description	Industry/Example	Key Benefits
Medical Imaging	Collaborative training for tumor detection and imaging analysis	NVIDIA Clara, RACOON, Owkin	Improved accuracy for rare diseases, privacy compliance
Diabetes & Chronic Disease Prediction	Risk modeling on EHR data across multiple hospitals	Multi-institutional diabetes studies	Early intervention, regulatory compliance
Oncology & Drug Discovery	Cancer prognosis and treatment response prediction	Owkin federated research networks	Accelerated research without data sharing
Infectious Disease Monitoring	Real-time pandemic modeling and surveillance	Public health federated systems (COVID-19)	Global collaboration, rapid response

These show that Federated Learning helps hospitals share data, especially where it's hard to do so normally. It speeds up research and improves finding diseases around the world (**Galtier & Owkin, 2023; Rieke et al., 2020; Tahir et al., 2025**).

Ethical Considerations

While Federated Learning is better for privacy, it also has ethical problems that need thinking about.

It helps protect data because it keeps patient info at each hospital and only shares model settings. This lowers the risk of big data leaks compared to putting everything in one place (**Rieke et al., 2020**).

But privacy isn't perfect. People can still try to figure out the data from the model updates. Also, attackers can put bad data in to mess up the model (**Tahir et al., 2025**). Methods like adding noise and secure ways to put updates together can help, but they cost some accuracy and time.

Bias is another problem. Hospitals have different patients, so the models can be unfair. For example, if only rich hospitals are involved, the model might not work well for poor people (**Pokhrel, 2024**). And if small hospitals can't join because they don't have enough resources, the models won't be as good for everyone.

Also, it's hard to know who's to blame when things go wrong with the models. And it takes a lot of computers and networks to run Federated Learning, which can keep some hospitals from joining.

Here's a table of the problems and solutions:

Ethical Issue	Description	Healthcare Impact	Mitigation Strategies
Privacy Vulnerabilities	Model inversion and poisoning attacks can compromise patient data	Potential patient data leaks and inaccurate diagnoses	Differential privacy, secure aggregation, robust validation
Algorithmic Bias & Fairness	Non-IID data from demographically imbalanced institutions	Disparate diagnostic performance across populations	Fairness-aware algorithms, diverse client inclusion
Accessibility & Equity	High computational requirements exclude resource-limited facilities	Widened healthcare gaps between regions	Lightweight FL variants, participation incentives
Accountability	Distributed system architecture complicates error attribution	Challenges in diagnostic error tracing and remediation	Comprehensive auditing protocols, governance frameworks

Addressing these multifaceted ethical challenges requires ongoing interdisciplinary collaboration

Fixing these problems needs people from different fields to work together to make rules that help Federated Learning do good while protecting people and being fair.

Conclusion

Federated Learning is a big step in making machine learning more private for healthcare. It lets hospitals work together without sharing patient info, which fixes privacy problems and keeps the models working well. Testing showed that it keeps over 97% of the accuracy of normal models.

This report shows that Federated Learning works well, protects privacy, and is used in things like medical images, predicting diseases, and watching for outbreaks. Companies like NVIDIA Clara and research groups are using it to help hospitals, drug companies, and health groups.

But using Federated Learning needs careful thinking about ethics. Problems like attacks, bias, and access need fixing with better privacy, fair algorithms, and ways to include everyone.

Looking ahead, we can make AI in healthcare even better by mixing Federated Learning with things like edge computing and blockchain. We can also make the models better for different data and make rules for everyone to follow. Also, making the AI easier to understand can help people trust it more.

Federated Learning helps make AI in healthcare fair and private, which helps improve health for everyone. Making this happen needs everyone to work together.

AI Tools Disclosure

I used help from computer systems at different times to do the research and writing:

- Grok (from xAI) was used a ton, to think of ideas, make an outline, find sources, and make some code to show how federated learning works.
- Grammarly was used at the end to find mistakes and make the writing sound better.

References

- Galtier, M., & Owkin, F. (2023). Applying AI to real-world health-care settings and the life sciences: Tackling data privacy, security and policy challenges with federated learning. *Artificial Intelligence in Science*, 171.
- Geyer, R. C., Klein, T., & Nabi, M. (2018). *Differentially Private Federated Learning: A Client Level Perspective* (arXiv:1712.07557). arXiv. <https://doi.org/10.48550/arXiv.1712.07557>
- Kumar, Y., Kaur, G. K., & Singh, R. (2024). A Survey on Diabetes Mellitus Prediction Using Machine Learning Methods. *2024 First International Conference on Technological Innovations and Advance Computing (TIACOMP)*, 391–396. <https://doi.org/10.1109/TIACOMP64125.2024.00072>
- McMahan, B., Moore, E., Ramage, D., Hampson, S., & Arcas, B. A. y. (2017). Communication-Efficient Learning of Deep Networks from Decentralized Data. *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, 1273–1282. <https://proceedings.mlr.press/v54/mcmahan17a.html>
- Pokhrel, A. (2024). Harmonizing public health with individual liberties: Exploring the interplay of right to health, privacy, and autonomy during recent and future pandemics. *International Journal of Human Rights in Healthcare*, 18(1), 107–111. <https://doi.org/10.1108/IJHRH-02-2024-0011>
- Rieke, N., Hancox, J., Li, W., Milletari, F., Roth, H. R., Albarqouni, S., Bakas, S., Galtier, M. N., Landman, B. A., & Maier-Hein, K. (2020). The future of digital health with federated learning. *NPJ Digital Medicine*, 3(1), 119.
- Tahir, N., Jung, C.-R., Lee, S.-D., Azizah, N., Ho, W.-C., & Li, T.-C. (2025). Federated Learning-Based Model for Predicting Mortality: Systematic Review and Meta-Analysis. *Journal of Medical Internet Research*, 27(1), e65708. <https://doi.org/10.2196/65708>